

Automatic Clone Recommendation for Refactoring Based on the Present and the Past

Ruru Yue^{*,†} Zhe Gao^{*,†} Na Meng[‡] Yingfei Xiong^{*,†} Xiaoyin Wang[§] J. David Morgenthaler[¶]

^{*}Key Laboratory of High Confidence Software Technologies (Peking University), MoE

[†]Institute of Software, EECS, Peking University, Beijing 100871, China
{yueruru,sonia,xiongyf}@pku.edu.cn

[‡]Department of Computer Science, Virginia Tech, Blacksburg, VA, 24060, nm8247@cs.vt.edu

[§]Department of Computer Science, University of Texas at San Antonio, San Antonio, TX 78249, xiaoyin.wang@utsa.edu

[¶]Google, Mountain View, CA 94043, jdm@google.com

Abstract—When many clones are detected in software programs, not all clones are equally important to developers. To help developers refactor code and improve software quality, various tools were built to recommend clone-removal refactorings based on the past and the present information, such as the cohesion degree of individual clones or the co-evolution relations of clone peers. The existence of these tools inspired us to build an approach that considers as many factors as possible to more accurately recommend clones. This paper introduces CREC, a learning-based approach that recommends clones by extracting features from the current status and past history of software projects. Given a set of software repositories, CREC first automatically extracts the clone groups historically refactored (R-clones) and those not refactored (NR-clones) to construct the training set. CREC extracts 34 features to characterize the content and evolution behaviors of individual clones, as well as the spatial, syntactical, and co-change relations of clone peers. With these features, CREC trains a classifier that recommends clones for refactoring.

We designed the largest feature set thus far for clone recommendation, and performed an evaluation on six large projects. The results show that our approach suggested refactorings with 83% and 76% F-scores in the within-project and cross-project settings. CREC significantly outperforms a state-of-the-art similar approach on our data set, with the latter one achieving 70% and 50% F-scores. We also compared the effectiveness of different factors and different learning algorithms.

I. INTRODUCTION

Code clones (or duplicated code) present challenges to software maintenance. They may require developers to repetitively apply similar edits to multiple program locations. When developers fail to consistently update clones, they may incompletely fix a bug or introduce new bugs [24], [33], [39], [42]. To mitigate the problem, programmers sometimes apply clone removal refactorings (e.g., *Extract Method* and *Form Template Method* [12]) to reduce code duplication and eliminate repetitive coding effort. However, as studied by existing work [15], [27], [51], not all clones need to be refactored. Many clones are not worthwhile to be refactored or difficult to be refactored [27], and programmers often only refactor a small portion of all clones [51]. Thus, when a clone detection tool reports lots of clones in a program, it would be difficult for the developer to go through the (usually long) list and locate the clones that should be refactored. A technique

that automatically recommends only the important clones for refactoring would be valuable.

Researchers have realized this problem and various tools were developed to recommend clones for refactoring. Given a project, most current tools recommend clones based on factors either representing *the present status or past evolution of the project*. For example, some tools suggest refactorings based on the code smell presented in a program, e.g., high cohesion and low coupling of clones [16], [19]. Some other tools provide recommendation based on the past co-evolution relationship between clones [18], [36]. The success of these tools indicates that both the present and the past of programs are useful in recommending clones, and it naturally raises a question: can we build an approach that considers as many factors as possible to more accurately recommend clones?

To take multiple factors into account, we need a model that aggregates all factors properly. Machine learning provides a good way to build a such model. Existing work [51] showed that using machine learning to integrate factors could lead to good accuracy in recommending clones. However, machine learning requires a large set of labelled instances for training. In practice, we often lack enough human power to manually build such a training set; neither have we found any existing work that automatically builds the set.

In this paper, we present CREC, a learning-based approach that automatically extracts refactored and non-refactored clones groups from software repositories, and trains models to recommend clones for refactoring. Specifically as shown in Fig. 1, there are three phases in our approach. Given open source project repositories, CREC first models clone genealogies for each repository as prior work does [27] and identifies clone groups historically refactored (**R-clones**) and others not (**NR-clones**) in Phase I.

CREC then extracts 34 features to characterize both kinds of groups and to train a model in Phase II. Intuitively, developers may decide refactorings based on (1) the cost of conducting refactorings (e.g., the divergence degree among clones), and (2) any potential benefit (e.g., better software maintainability). To consider as many factors as possible, we thus include 17 per-clone features to characterize each clone's content and evolution history, and 17 clone-group features

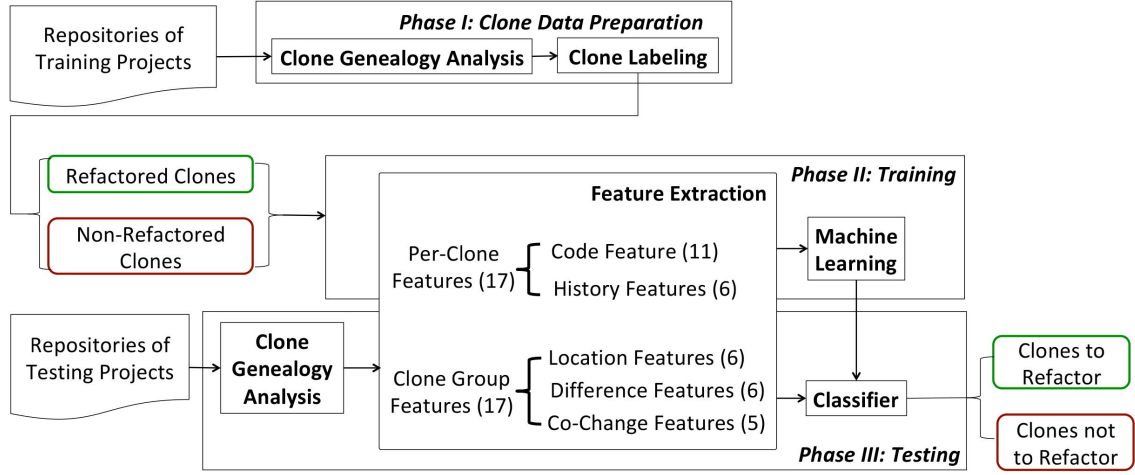


Fig. 1: CREC has three phases: clone data preparation, training, and testing. Phase I detects refactored and non-refactored clones in the repositories of training projects. Phase II trains a classifier with machine learning based on the two types of clones. Phase III recommends clones for refactoring using the trained classifier.

to capture the relative location, syntactic difference, and co-change relationship between clone peers. The 34 features in total describe both the present status and past evolution of single clones and clone peers.

Finally in Phase III, to recommend clones given a project’s repository, CREC performs clone genealogy analysis, extracts features from the clone genealogies, and uses the trained model for refactoring-decision prediction.

We evaluated CREC with R-clones and NR-clones extracted from six open source projects in the within-project and cross-project settings. Based on our dataset and extracted features, we also evaluated the effectiveness of different feature subsets and machine learning algorithms for clone refactoring recommendation. Our evaluation has several findings. (1) CREC effectively suggested clones for refactoring with 83% F-score in the within-project setting and 76% F-score in the cross-project setting. (2) CREC significantly outperforms a state-of-the-art learning-based approach [51], which has 70% F-score in the within-project setting and 50% F-score in the cross-project setting. (3) Our automatically extracted refactored clones achieve a precision of 92% on average. (4) Clone history features and co-change features are the most effective features, while tree-based machine learning algorithms are more effective than SMO and Naive Bayes.

In summary, we have made the following contributions:

- We developed CREC, a novel approach that (1) automates the extraction of R-clones from repositories in an accurate way, and (2) recommends clones for refactoring based on a comprehensive set of characteristic features.
- We designed 34 features—the largest feature set thus far—to characterize the content and evolution of individual clones, as well as the spatial, syntactical, and co-evolution relationship among clones.
- We conducted a systematic evaluation to assess (1) the precision of automatic R-clone extraction, (2) CREC’s effectiveness of clone recommendation in comparison with a state-of-the-art approach [51], and (3) the tool’s

sensitivity to the usage of different feature sets and various machine learning algorithms.

II. APPROACH

CREC consists of three phases. In this section, we will first discuss how CREC automatically extracts R-clones and NR-clones from software history (Section II-A). Then we will explain our feature extraction to characterize clone groups (Section II-B). Finally, we will expound on the training and testing phases (Section II-C and II-D).

A. Clone Data Preparation

This phase contains two steps. First, CREC conducts clone genealogy analysis to identify clone groups and extract their evolution history (Section II-A1). Second, based on the established clone genealogies, CREC implements several heuristics to decide whether a clone group was refactored or not by developers (Section II-A2).

1) *Clone Genealogy Analysis*: Clone genealogy describes how clones evolve over time. Kim et al. invented an approach to model clone genealogies based on software version history [27]. As the original approach implementation is not available, we reimplemented the approach with some improvement on multi-version clone detection.




Fig. 2: An exemplar clone lineage

Fig. 2 illustrates a clone lineage modeled by our analysis. Suppose there are three clone groups detected in a project’s

different versions: $G_1 = \{A, B, C, D\}$, $G_2 = \{A', B', C'\}$, and $G_3 = \{A'', B''\}$. By comparing clones' string similarity across versions, we learnt that G_1 and G_2 have three similar but different matching clone members: (A, A') , (B, B') , and (C, C') ; while G_2 and G_3 have two matching members: (A', A'') and (B', B'') . Therefore, we correlate the three clone groups. Furthermore, we infer that G_1 evolves to G_2 because the majority of their members match. G_2 evolves to G_3 , because two members seem to consistently change with the third member (C') unchanged. Therefore, there are mainly two parts in our clone genealogy analysis: multi-version clone detection and cross-version clone linking.

Multi-Version Clone Detection. We utilized SourcererCC [44], a recent scalable and accurate near-miss clone detector, to find block-level clones. As with prior work [43], [44], we configured SourcererCC to report clones with at least 30 tokens or 6 lines of code to avoid trivial code duplication. We also set SourcererCC's minimum similarity threshold as 80% to reduce false alarms.

Since the whole history may contain a huge number of commits, to efficiently establish clone genealogies, we sampled versions in chronological order based on the amount of code revision between versions. Specifically, CREC samples the oldest version of a project by default, denoting it as S_1 . Then CREC enumerates subsequent versions until finding one that has at least 200 changed lines in comparison with S_1 , denoting it as S_2 . By selecting S_2 as another sample, CREC further traverses the follow-up versions until finding another version that is different from S_2 by at least 200 lines of code. This process continues until CREC arrives at the latest version.

Cross-Version Clone Linking. We used an existing approach [52] to link the clone groups detected in different sampled versions, in order to restore the clone evolution history. This approach matches clones across versions based on the code location and textual similarity. For example, suppose that two clones A and B are found in file F of version V_i ; while one clone A' is also found in F , but of version V_{i+1} . Since A and A' are located in the same file, and if A has more content overlap with A' than B does, we link A with A' across versions, and conclude that A evolves to A' . With such clone linking, we can further link clone groups across versions based on the membership of linked clones.

2) *Clone Labeling*: To train a classifier for clone refactoring recommendation, we need to label clones either as "refactored" or "not refactored" by developers. Intuitively, if (1) certain clones have duplicated code removed and method invocations added, and (2) the newly invoked method has content similar to the removed code, we infer that developers might have applied "Extract Method" refactorings to these clones for duplication removal. Based on this intuition, we developed CREC to assess any given clone group G of version V_i with the following three criteria:

- i. At least two clone instances in G have code size reduced when evolving into clones of version V_{i+1} . We refer to this subset of clones as $C = \{c_1, c_2, \dots\} (C \subset G)$.

- ii. Within C , at least two clone instances (from the same or different source files) add invocations to a method $m()$ during the evolution. We refer to the corresponding clone subset as $C' = \{c'_1, c'_2, \dots\} (C' \subset C)$.
- iii. Within C' , at least two clone members have their removed code similar to $m()$'s body. Here, we set the similarity threshold as 0.4, because the implementation of $m()$ may extract the commonality of multiple clones while parameterizing any differences, and thus become dissimilar to each original clone. This threshold setting will be further discussed in Section III-C.

When G satisfies all above criteria, CREC labels it as "R-clone"; otherwise, it is labeled as "NR-clone".

B. Feature Extraction

With the ability to detect so many clones, determining which of them to refactor can be challenging. Prior research shows that developers base their refactoring decisions on the potential efforts and benefits of conducting refactorings [28], [51]. Our intuition is that by extracting as many features as possible to reflect the present status and history of programs, we can capture the factors that may affect developers' refactoring decisions, and thus accurately suggest clones for refactoring. Specifically, we extract five categories of features: two categories modeling the current version and historic evolution of individual clones, and three categories reflecting the location, code difference, and co-change relationship among clone peers of each group. Table I summarizes the 34 features of these 5 categories.

1) *Clone Code Features*: We believe that the benefits and efforts of refactoring a clone are relevant to the code size, complexity, and the clone's dependence relationship with surrounding code. For instance, if a clone is large, complex, and independent of its surrounding context, the clone is likely to be refactored. Therefore, we defined 11 features (F1-F11) to characterize the content of individual clones.

2) *Clone History Features*: A project's evolution history may imply its future evolution. Prior work shows that many refactorings are driven by maintenance tasks such as feature extensions or bug fixes [45]. It is possible that developers refactor a clone because the code has recently been modified repetitively. To model a project's recent evolution history, we took the latest 1/10 of the project's sampled commits, and extracted 6 features (F12-F17).

3) *Relative Location Features among Clones*: Clones located in the same or closely related program components (e.g., packages, files, or classes) can share program context in common, and thus may be easier to refactor for developers. Consequently, we defined 6 features (F18-F23) to model the relative location information among clones, characterizing the potential refactoring effort.

4) *Syntactic Difference Features among Clones*: Intuitively, dissimilar clone instances are usually more difficult to refactor than similar ones, because developers have to resolve the differences before extracting common code as a new method. To estimate the potential refactoring effort, we defined 6

TABLE I: The 34 features used in CREC

| Feature Category | Features | Rationale |
|--|---|---|
| Clone Code (11) | F1: Lines of code in a clone. | A larger clone is more likely to be refactored to significantly reduce code size. |
| | F2: Number of tokens. | Similar to F1. |
| | F3: Cyclomatic complexity. | The more complexity a clone has, the harder it is to maintain the clone, and thus the more probably developers will refactor. |
| | F4: Number of field accessed by a clone. | The more fields (e.g., private fields) a clone accesses, the more this clone may depend on its surrounding program context. Consequently, refactoring this clone may require more effort to handle the contextual dependencies. |
| | F5: Percentage of method invocation statements among all statements in a clone. | Similar to F4, this feature also measures the dependence a clone may have on its program context. |
| | F6: Percentage of code lines of a clone out of its container method. | Similar to F4 and F5, this feature indicates the context information of a clone in its surrounding method and how dependent the clone may be on the context. |
| | F7: Percentage of arithmetic statements among all statements in a clone. | Arithmetic operations (e.g., “+” and “/”) may indicate intensive computation, and developers may want to refactor such code to reuse the computation. |
| | F8: Whether a complete control flow block is contained. | If a clone contains an incomplete control flow, such as an <code>else-if</code> branch, it is unlikely that developers will refactor the code. |
| | F9: Whether a clone begins with a control statement. | Beginning with a control statement (e.g., <code>for</code> -loop) may indicate separable program logic, which can be easily extracted as a new method. |
| | F10: Whether a clone follows a control statement. | Similar to F9, a control statement may indicate separable logic. |
| | F11: Whether a clone is test code. | Developers may intentionally create clones in test methods to test similar or relevant functionalities, so they may not want to remove the clones. |
| Clone History (6) | F12: Percentage of file-existence commits among all checked historical commits. | If a file is recently created, developers may not refactor any of its code. |
| | F13: Percentage of file-change commits among all checked commits. | If a file has been changed a lot, developers may want to refactor the code to facilitate software modification. |
| | F14: Percentage of directory-change commits among all checked commits. | Similar to F13, this feature also models historical change activities. |
| | F15: Percentage of recent file-change commits among the most recent 1/4 checked commits. | Compared with F13, this feature focuses on a shorter history, because a file’s most recent updates may influence its future more than the old ones. |
| | F16: Percentage of recent directory-change commits among the most recent 1/4 checked commits. | Compared with F14, this feature also measures how actively a directory was modified, but characterizes a shorter history. |
| | F17: Percentage of a file’s maintainers among all developers. | The more developers modify a file, the more likely that the file is buggy [37]; and thus, developers may want to refactor such files. |
| Relative Locations among Clones (6) | F18: Whether clone peers in a group are in the same directory | It seems easier to refactor clones within the same directory, because clones share the directory as their common context. |
| | F19: Whether the clone peers are in the same file. | Similar to F18, clones’ co-existence in the same file may also indicate less refactoring effort. |
| | F20: Whether the clone peers are in the same class hierarchy | Similar to F18 and F19, if clones share their accesses to many fields and methods defined in the same hierarchy, such clones may be easier to refactor. |
| | F21: Whether the clone peers are in the same method | Similar to F18-F20, clones within the same method possibly share many local variables, indicating less effort to apply refactorings. |
| | F22: Distance of method names measures the minimal Levenshtein distance [32] (i.e., textual edit distance) among the names of methods containing clone peers. | As developers define method names as functionality descriptors, similar method names usually indicate similar implementation, and thus may imply less refactoring effort. |
| | F23: Possibility of the clone peers from intentionally copied directories. | The more similar two directories’ layout are to each other, the more likely that developers intentionally created the cloned directories, and the less likely that clones between these directories are to be refactored. |
| Syntactic Differences among Clones (6) | F24: Number of instances in a clone group | The size of a clone group can affect developers’ refactoring decision. For instance, refactoring a larger group may lead to more benefit (e.g., reduce code size), while refactoring a smaller group may require for less effort. |
| | F25: Number of total diffs measures the number of differential multisets reported by MCIDiff [34]. | More differences may demand for more refactoring effort. |
| | F26: Percentage of partially-same diffs measures among all reported differential multisets, how many multisets are partially same. | Partially-same multisets indicate that some instances in a clone group are identical to each other while other instances are different. Developers may easily refactor the clone subset consisting of identical clones without dealing with the difference(s) in other clone peers. |
| | F27: Percentage of variable diffs measures among all differential multisets, how many multisets contain variable identifiers. | If developers resolve each variable diff by renaming the corresponding variables consistently or declaring a parameter in the extracted method, this feature characterizes such effort. |
| | F28: Percentage of method diffs measures among all differential multisets, how many multisets have method identifiers. | If developers resolve each method diff by creating a lambda expression or constructing template methods, this feature reflects the amount of such effort. |
| | F29: Percentage of type diffs measures among all differential multisets, what is the percentage of multisets containing type tokens. | Different from how variable and method differences are processed, type differences are possibly resolved by generalizing the types into a common super type. This feature reflects such potential refactoring effort. |
| | F30: Percentage of all-change commits among all checked commits. | The more often all clones are changed together, the more possibly developers will refactor the code. |
| | F31: Percentage of no-change commits among all checked commits. | Similar to F30, this feature also characterizes the evolution similarity among clone peers. |
| Co-Changes among Clones (5) | F32: Percentage of 1-clone-change commit among all checked commits. | Different from F30, this feature characterize the inconsistency evolution of clones. If clones divergently evolve, they are less likely to be refactored. |
| | F33: Percentage of 2-clone-change commits among all checked commits. | Similar to F32, this feature also captures the maintenance inconsistency among clones. |
| | F34: Percentage of 3-clone-change commits among all checked commits. | Similar to F32 and F33. |

features (F24-F29) to model differences among clones, such as the divergent identifier usage of variables, methods, and types.

Specifically to extract differences among clones, we used an off-the-shelf tool—MCIDiff [34]. Given a clone group $G = \{c_1, c_2, \dots, c_n\}$, MCIDiff tokenizes each clone, leverages the longest common subsequence (LCS) algorithm to match token sequences, and creates a list of matched and differential multisets of tokens. Intuitively, if a token commonly exists in all clones, it is put into a **matched** multiset; otherwise, it is put to a **differential** multiset. Furthermore, there is a special type of differential multisets called **partially same** to record tokens that commonly exist in some (≥ 2) but not all clones in a group. Matched and partially-same multisets may imply zero refactoring effort, because developers do not need to resolve differences before refactoring the tokens’ corresponding code.

Meng et al. showed that variable differences among clones were easier to resolve than method or type differences [36]. To more precisely characterize the potential refactoring effort, we also classified identifier tokens into three types: variables, methods, and types. As MCIDiff does not identify different types of tokens, we extended it to further analyze the syntactic meaning of identifiers. With more detail, our extension first parsed the Abstract Syntax Tree (AST) of each clone using Eclipse ASTParser [1]. By implementing an ASTVisitor to traverse certain types of ASTNodes, such as `FieldAccess`, `MethodInvocation`, and `TypeLiteral`, we determined whether certain token(s) used in these ASTNodes are names of variables, methods, or types. In this way, we further classified identifier tokens and labeled MCIDiff’s differential multisets with “variable”, “method”, and “type” accordingly.

5) *Co-Change Features among Clones*: If multiple clone peers of the same group are frequently changed together, developers may want to refactor the code to reduce or eliminate future repetitive editing effort [35]. Thus, we defined 5 features (F30-F34) to reflect the co-change relationship among clones.

C. Training

To train a binary-class classifier for clone recommendation, we first extracted feature vectors for both R-clones and NR-clones in the training set. Then for each feature vector we constructed one data point in the following format: $\langle feature_vector, label \rangle$, where *feature_vector* consists of 34 features; and *label* is set to 1 or 0 depending on whether *feature_vector* is derived from R-clones or NR-clones.

We used Weka [2]—a software library implementing a collection of machine learning algorithms to train a classifier based on training data. By default, we use AdaBoost [13], a state-of-the-art machine learning algorithm, because prior research shows that it can be successfully used in many applications [11]. Specifically, we configured AdaBoost to take a weak learning algorithm, decision stump [22], as its weak learner, and call the weak learner 50 times.

D. Testing

Given the clone groups in the testing set, CREC first extracts features based on the current status and evolution

history of these clones. Next, CREC feeds the features to the trained classifier to decide whether the clone groups should be refactored or not. If a group is predicted to be refactored with over 50% likelihood, CREC recommends the group for refactoring. When multiple groups are recommended, CREC further ranks them based on the likelihood outputs by the classifier.

III. EXPERIMENTS

In this section, we first introduce our data set for evaluation (Section III-A), and then describe our evaluation methodology (Section III-B). Based on the methodology, we estimated the precision of reporting refactored clone groups (Section III-C). Next, we evaluated the effectiveness of CREC and compare it with an existing clone recommendation approach (Section III-D). Finally, we identified important feature subsets (Section III-E) and suitable machine learning algorithms (Section III-F) for clone refactoring recommendation.

A. Data Set Construction

Our experiments used the repositories of six open source projects: Axis2, Eclipse.jdt.core, Elastic Search, JFreeChart, JRuby, and Lucene. We chose these projects for two reasons. First, they have long version histories, and most of these projects were also used by prior research [36], [46], [49]–[51]. Second, these projects belong to different application domains, which makes our evaluation results representative. Table II presents relevant information of these projects. Column **Domain** describes each project’s domain. **Study Period** is the evolution period we investigated. **# of Commits** is the total number of commits during the studied period. **# of Clone Groups** is the total number of clone groups detected in each project’s sampled versions. **# of R-clones** means the number of refactored clone groups automatically extracted by CREC.

To evaluate CREC’s effectiveness, we need both R-clones and NR-clones to construct the training and testing data sets. As shown in Table II, there are a small portion of detected clone groups reported as R-clones. To create a balanced data set of positive and negative examples, we randomly chose a subset of clone groups from clone genealogies without R-clones. The number of NR-clones in the subset is equal to that of reported R-clones. Each time when we trained a classifier, we used a portion of this balanced data set as **training data**. For **testing data**, we used the remaining portion of the data set as the testing set, except the R-clones’ false positives confirmed by our manual inspection (See Section III-C). We did not manually inspect the NR-clones in the testing set. This is because existing study [51] has revealed only a very small portion of clone groups (less than 1%) are actually refactored by developers. Therefore, the portion of false NR-clones (i.e., missing true R-clones) in the testing set cannot exceed 1%, which would have only marginal effect on our results.

B. Evaluation Methodology

CREC automatically extracts clone groups historically refactored (R-clones) from software repositories, and considers the

TABLE II: Subject projects

| Project | Domain | Study Period | # of Commits | # of Clone Groups | # of R-clones |
|------------------|--|-----------------------|--------------|-------------------|---------------|
| Axis2 | core engine for web services | 09/02/2004-03/06/2018 | 8723 | 800950 | 43 |
| Eclipse.jdt.core | integrated development environment | 06/05/2001-01/22/2018 | 22358 | 10001106 | 106 |
| Elastic Search | search engine | 02/08/2010-01/05/2018 | 14766 | 1856035 | 33 |
| JFreeChart | open-source framework for charts | 06/29/2007-01/24/2018 | 3603 | 246098 | 59 |
| JRuby | an implementation of the ruby language | 09/10/2001-03/04/2018 | 24434 | 718698 | 65 |
| Lucene | information retrieval library | 09/11/2001-01/06/2018 | 22061 | 4774413 | 27 |

remaining clone groups as not refactored (NR-clones). Before using the extracted groups to construct the training set, we first assessed CREC’s precision of R-clone extraction.

Additionally, with the six project’s clone data, we conducted within-project and cross-project evaluation. In *within-project evaluation*, we used ten-fold validation to evaluate recommendation effectiveness. In *cross-project evaluation*, we iteratively used five projects’ R-clones and NR-clones as training data, and used the remaining project’s clone data for testing. In both cases, the clone recommendation effectiveness is evaluated with three metrics: precision, recall, and F-score.

Precision (P) measures among all the recommended clone groups for refactoring, how many groups were actually refactored by developers:

$$P = \frac{\text{\# of refactored clone groups}}{\text{Total \# of recommended groups}} * 100\%. \quad (1)$$

Recall (R) measures among all known refactored clone groups, how many groups are suggested by a refactoring recommendation approach:

$$R = \frac{\text{\# of refactored clone groups suggested}}{\text{Total \# of known refactored groups}} * 100\%. \quad (2)$$

F-score (F) is the harmonic mean of precision and recall. It combines P and R to measure the overall accuracy of refactoring recommendation as below:

$$F = \frac{2 * P * R}{P + R} * 100\%. \quad (3)$$

Precision, recall, and F-score all vary within [0%, 100%]. The higher precision values are desirable, because they demonstrate better accuracy of refactoring recommendation. As the within-project and cross-project evaluation were performed with six projects, we obtained six precision, recall, and F-score values respectively, and treated the averages of all six projects as the overall effectiveness measurement.

C. Precision of Reporting R-clones

CREC relies on a similarity threshold parameter l_{th} to determine R-clones (see Section II-A2). With the threshold, CREC checks whether the implementation of a newly invoked method is sufficiently similar to a removed code snippet. In our experiment, we changed l_{th} from 0.3 to 0.5 with 0.1 increment, and manually examined the reported R-clones in the six projects to estimate CREC’s precision.

In Table III, **# of Reported** shows the total number of reported R-clones by CREC. **# of Confirmed** is the number

of reported R-clones manually confirmed. **Precision** measures CREC’s precision. As l_{th} ’s value increases, the total number of reported groups goes down, while the precision goes up. To achieve a good balance between the total number of reported refactored groups and the precision, we set l_{th} to 0.4 by default. This is because when $l_{th} = 0.4$, the number of R-clones reported by CREC and the 92% precision are sufficient for constructing the training set.

TABLE III: CREC’s precision of reporting R-clones

| Threshold Value | # of Reported | # of Confirmed | Precision |
|-----------------|---------------|----------------|-----------|
| 0.3 | 577 | 460 | 80% |
| 0.4 | 333 | 307 | 92% |
| 0.5 | 204 | 193 | 95% |

D. Effectiveness Comparison of Clone Recommendation Approaches

Wang et al. developed a machine learning-based clone recommendation approach [51], which extracted 15 features for cloned snippets and leveraged C4.5 to train and test a classifier. Among the 15 features, 10 features represent the current status of individual clones, 1 feature reflects the evolution history of individual clones, 3 features describe the spatial relationship between clones, and 1 feature reflects any syntactic difference between clones. Wang et al.’s feature set is much smaller than our 34-feature set. It fully overlaps with our feature category of clone content, and matches three out of our six spatial location-related features. Wang et al.’s feature set characterizes individual clones more comprehensively than the relationship between clone peers (11 features vs. 4 features), and emphasizes clones’ present versions more than their previous versions (14 features vs. 1 feature). It does not include any feature to model the co-change relationship between clones.

In this experiment, we compared CREC with Wang et al.’s approach to see how CREC performs differently. Because neither the tool nor data of Wang et al.’s paper is accessible even after we contacted the authors, we reimplemented Wang et al.’s approach, and used our data set for evaluation. As CREC is different from Wang et al.’s approach in terms of both the feature set and machine learning (ML) algorithm, we also developed two variants of our approach: CREC_m that uses CREC’s ML algorithm (AdaBoost) but the feature set in Wang et al.’s work, and CREC_f that uses CREC’s feature set but the ML algorithm (C4.5) in Wang et al.’s work. Next, we evaluated the four clone recommendation approaches with the

six projects’ data in both the within-project and cross-project settings. Note that due to different data sets for evaluation, the effectiveness of Wang et al.’s approach in this paper is different from the results reported in their paper.

Table IV presents the four approaches’ effectiveness comparison for within-project prediction. On average, CREC achieved 82% precision, 86% recall, and 83% F-score, which were significantly higher than the 73% precision, 71% recall, and 70% F-score by Wang et al.’s approach. CREC improved the overall average F-score of Wang et al.’s approach (70%) by 19%. This implies that compared with Wang et al.’s approach, CREC more effectively recommended clones for refactoring.

To further analyze why CREC worked differently from Wang et al.’s approach, we also compared their approach with CREC_m and CREC_f. As shown in Table IV, the F-score of CREC_m was a little higher than Wang et al.’s approach while CREC_f worked far more effectively than both of them. The comparison indicates that (1) both our feature set and AdaBoost positively contribute to CREC’s better effectiveness over Wang et al.’s approach; and (2) our feature set can more significantly improve Wang et al.’s approach than AdaBoost.

Table V shows different approaches’ evaluation results for cross-project prediction. CREC achieved 77% precision, 77% recall, and 76% F-score on average. In comparison, Wang et al.’s approach acquired 53% precision, 50% recall, and 50% F-score. By replacing their feature set with ours, we significantly increased the precision to 72%. We also replaced their C4.5 algorithm with the AdaBoost algorithm used in CREC, and observed that the F-score increased from 50% to 53%. By correlating these observations with the above-mentioned observations for the within-project setting, we see that our 34-feature set always works much better than Wang et al.’s 15-feature set. Additionally, the replacement of C4.5 with AdaBoost brings small improvement to the effectiveness of Wang et al.’s approach. Therefore, the rationale behind our feature set plays a crucially important role in CREC’s significant improvement over the state-of-the-art approach.

Finding 1: CREC outperforms Wang et al.’s approach, mainly because it leverages 34 features to comprehensively characterize the present and past of code clones.

On average, CREC achieved 82% precision, 86% recall, and 83% F-score in the within-project setting, which were higher than the 77% precision, 77% recall, and 76% F-score in the cross-project setting. This observation is as expected, because the clones of one project may share certain project-specific commonality. By training and testing a classifier with clone data of the same project, CREC can learn some project-specific information, which allows CREC to effectively suggest clones.

CREC worked better for some projects but worse for the others. For instance, CREC obtained the highest F-score (95%) for Lucene and the lowest F-score (66%) for Elastic Search in the within-project setting. This F-score value range means that for the majority of cases, CREC’s recommendation is meaningful to developers. It implies two insights. First, developers’

refactoring decisions are generally predictable, which validates our hypothesis that different developers similarly decide what clones to refactor. Second, some developers may refactor clones more frequently than the others and have personal refactoring preferences. Such personal variance can cause the prediction divergence among different test projects.

Finding 2: CREC suggests refactorings with 66-95% F-scores within projects and 54-93% F-scores across projects. It indicates that developers’ clone refactoring decisions are generally predictable, although the refactoring preference may vary with people and projects.

E. Important Feature Subset Identification

CREC leveraged 5 categories of features, including 11 features to characterize the content of individual clones, 6 features to reflect individual clones’ evolution history, 6 features to represent the spatial locations of clone peers, 6 features to capture any syntactic difference between clones, and 5 features to reflect the co-evolution relationship between clones. To explore how each group of features contribute to CREC’s overall effectiveness, we also created five variant feature sets by removing one feature category from the original set at a time. We then executed CREC with each variant feature set in both the within-project and cross-project settings (see Fig. 3).

In the within-project setting, CREC’s all-feature set achieved 83% F-score. Interestingly, **ExceptDiff** obtained the same F-score, which means that the features capturing syntactic differences between clone peers were not unique to characterize R-clones. This category of features was inferable from other features, and thus removing this category did not influence CREC’s effectiveness. **ExceptCode** and **ExceptLocation** obtained the highest F-score—84%. This may imply that the two feature subsets negatively affected CREC’s effectiveness, probably because the code smell and the location information in clones were the least important factors that developers considered when refactoring code. **ExceptHistory** obtained the lowest F-score—77% and **ExceptCoChange** achieved the second lowest F-score—81%. This indicates that the evolution history of individual clones and the co-change relationship of clone peers provide good indicators for clone refactorings. There seems to be a strong positive correlation between developers’ refactoring decisions and clones’ evolution history.

In the cross-project setting, CREC’s all feature set achieved 76% F-score. **ExceptDiff** obtained exactly the same F-score, probably because the features capturing any syntactic differences between clones are inferable from other features. All other variant feature sets led to lower F-scores. It means that except the syntactic differences between clones, the other four categories of features positively contribute to CREC’s prediction performance. Especially, **ExceptHistory** obtained an F-score much lower than the scores of other feature sets. This phenomenon coincides with our observations in the within-project setting, implying that the evolution history of

TABLE IV: Effectiveness comparison of different clone recommendation approaches in the within-project setting

| Test Project | CREC (CREC's Feature Set+AdaBoost) | | | CREC _m (Wang et al.'s Feature Set+AdaBoost) | | | CREC _f (CREC's Feature Set+C4.5) | | | Wang et al.'s Approach (Wang et al.'s Feature Set+C4.5) | | |
|------------------|---------------------------------------|-------|-------|---|-------|-------|--|-------|-------|--|-------|-------|
| | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) |
| Axis2 | 79 | 86 | 81 | 80 | 70 | 73 | 74 | 88 | 79 | 76 | 64 | 65 |
| Eclipse.jdt.core | 85 | 84 | 83 | 81 | 79 | 79 | 83 | 84 | 82 | 77 | 79 | 77 |
| Elastic Search | 64 | 70 | 66 | 52 | 55 | 50 | 69 | 80 | 74 | 46 | 45 | 45 |
| JFreeChart | 91 | 95 | 92 | 78 | 78 | 78 | 90 | 98 | 94 | 76 | 83 | 77 |
| JRuby | 79 | 87 | 82 | 66 | 74 | 69 | 83 | 81 | 82 | 69 | 70 | 69 |
| Lucene | 95 | 97 | 95 | 77 | 83 | 79 | 93 | 97 | 94 | 93 | 87 | 87 |
| Average | 82 | 86 | 83 | 72 | 73 | 71 | 82 | 88 | 84 | 73 | 71 | 70 |

TABLE V: Effectiveness comparison of different clone recommendation approaches in the cross-project setting

| Test Project | CREC (CREC's Feature Set+AdaBoost) | | | CREC _m (Wang et al.'s Feature Set+AdaBoost) | | | CREC _f (CREC's Feature Set+C4.5) | | | Wang et al.'s Approach (Wang et al.'s Feature Set+C4.5) | | |
|------------------|---------------------------------------|-------|-------|---|-------|-------|--|-------|-------|--|-------|-------|
| | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) | P (%) | R (%) | F (%) |
| Axis2 | 82 | 84 | 83 | 62 | 42 | 50 | 80 | 53 | 64 | 55 | 32 | 40 |
| Eclipse.jdt.core | 71 | 90 | 80 | 57 | 79 | 66 | 63 | 86 | 73 | 67 | 83 | 74 |
| Elastic Search | 70 | 44 | 54 | 55 | 50 | 53 | 67 | 31 | 43 | 48 | 38 | 42 |
| JFreeChart | 73 | 86 | 79 | 39 | 25 | 31 | 46 | 20 | 28 | 33 | 25 | 29 |
| JRuby | 72 | 63 | 67 | 60 | 56 | 58 | 75 | 83 | 78 | 66 | 77 | 71 |
| Lucene | 93 | 93 | 93 | 57 | 63 | 60 | 100 | 48 | 65 | 48 | 44 | 46 |
| Average | 77 | 77 | 76 | 55 | 53 | 53 | 72 | 54 | 58 | 53 | 50 | 50 |




Fig. 3: CREC's average effectiveness when using different feature sets

code clones are always positive indicators of clone refactorings. **ExceptCoChange** acquired the second lowest F-score, meaning that the co-evolution relationship of clone peers is also important to predict developers' refactoring practice.

Finding 3: The clone history features of single clones and co-change features among clone peers were the two most important feature subsets. It seems that developers relied more on clones' past than their present to decide refactorings.

F. Suitable Model Identification

To understand which machine learning (ML) algorithms are suitable for CREC, in addition to AdaBoost, we also experimented with four other supervised learning algorithms: Random Forest [20], C4.5 [41], SMO [40], and Naive Bayes [3].

All these algorithms were executed with their default parameter settings and were evaluated on the six projects for within-project and cross-project predictions. Fig. 4 presents our experiment results. The three bars for each ML algorithm in each chart separately illustrate the average precision, recall, and F-score among six projects.

In the within-project setting, Random Forest and C4.5 achieved the highest F-score: 84%. AdaBoost obtained the second best F-score: 83%. The F-score of Naive Bayes was the lowest—64%, and SMO derived the second lowest F-score: 78%. In the cross-project setting, AdaBoost's F-score (76%) was much higher than the F-scores of other algorithms. Random Forest obtained the second highest F-score: 61%. C4.5 acquired the third highest F-score—58%, while SMO and Naive Bayes had the lowest F-score: 53%. Between the two settings, AdaBoost reliably performed well, while Naive Bayes was always the worst ML algorithm for clone refactoring




Fig. 4: CREC's average effectiveness when using different ML algorithms

recommendation. AdaBoost, Random Forest and C4.5 are all tree-based ML algorithms, while SMO and Naive Bayes are not. The results in Fig. 4 imply that tree-based ML algorithms are more suitable for clone recommendation than others.

Finding 4: AdaBoost reliably suggests clones for refactoring with high accuracy in both within-project and cross-project settings, while Naive Bayes consistently worked poorly. Among the five experimented ML algorithms, tree-based algorithms stably worked better than the other algorithms.

IV. RELATED WORK

The related work includes clone evolution analysis, empirical studies on clone removals, clone recommendation for refactoring, and automatic clone removal refactoring.

A. Clone Evolution Analysis

Kim et al. initiated clone genealogy analysis and built a clone genealogy extractor [27]. They observed that extensive refactorings of short-lived clones may not be worthwhile, while some long-lived clones are not easily refactorable.

Other researchers conducted similar clone analysis and observed controversial or complementary findings [4], [10], [15], [30], [47]. For instance, Aversano et al. found that most co-changed clones were consistently maintained [4], while Krinke observed that half of the changes to clones were inconsistent [30]. Göde et al. discovered that 88% of clones were never changed or changed only once during their lifetime, and only 15% of all changes to clones were unintentionally inconsistent [15]. Thummalapenta et al. automatically classified clone evolution into three patterns: consistent evolution, independent evolution, and late propagation [47]. Cai et al. analyzed long-lived clones, and reported that the number of developers who modified clones and the time since the last addition or removal of a clone to its group were highly correlated with the survival time of clones [10].

These studies showed that not every detected clone should be refactored, which motivated our research that recommends only clones likely to be refactored.

B. Empirical Studies on Clone Removals

Various studies were conducted to understand how developers removed clones [14]. Specifically, Göde observed that developers deliberately removed clones with method extraction refactorings, and the removed clones were mostly co-located in the same file [14]. Bazrafshan et al. extended Göde's study, and found that the accidental removals of code clones occurred slightly more often than deliberate removals [8]. Similarly, Choi et al. also observed that most clones were removed with either *Extract Method* or *Replace Method with Method* refactorings. Silva et al. automatically detected applied refactorings in open source projects on GitHub, and sent questionnaires to developers for the motivation behind their refactoring practice [45]. They found that refactoring activities were mainly driven by changes in the requirements, and *Extract Method* was the most versatile refactoring operation serving 11 different purposes. None of these studies automatically recommends clones for refactoring or ranks the recommendations.

C. Clone Recommendation for Refactoring

Based on code clones detected by various techniques [23], [26], [29], a variety of tools identify or rank refactoring opportunities [5], [7], [16], [18], [19], [48]. For instance, Balazinska et al. defined a clone classification scheme based on various types of differences among clones, and automated the classification to help developers assess refactoring opportunities for each clone group [5]. Balazinska et al. later built another tool that facilitated developers' refactoring decisions by presenting the syntactic differences and common contextual dependencies among clones [7]. Higo et al. and Goto et al. ranked clones based on the coupling or cohesion metrics [16], [19]. Tsantalidis et al. and Higo et al. further improved the ranking mechanism by prioritizing clones that had been repetitively changed in the past [18], [48].

Wang et al. extracted 15 features to characterize the content and history of individual clones, and the context and difference relationship among clones [51]. Replacing the feature set with CREC's boost the overall effectiveness of Wang et al.'s approach. More importantly, our research shows for the first time that history-based features are more important than smell-

based features (i.e., clone content, location, or differences) when suggesting relevant clones for refactoring.

D. Automatic Clone Removal Refactoring

A number of techniques automatically remove clones by applying refactorings [6], [9], [17], [21], [25], [31], [36], [46], [49], [50]. These tools mainly extract a method by factorizing the common part and parameterizing any differences among clones. For example, Tairas et al. built an Eclipse plugin called CeDAR, which unifies clone detection, analysis and refactoring [46]. Juillerat et al. and Balazinska et al. defined extra mechanisms to mask any difference between clones [6], [25]. Several researchers conducted program dependency analysis to further automate *Extract Method* refactorings for near-miss clones—clones that contain divergent program syntactic structures or different statements [9], [21], [31], [31]. Specifically, Hotta et al. safely moved out the distinct statements standing between duplicated code to uniformly extract common code [21]. Krishnan et al. identified the maximum common subgraph between two clones' Program Dependency Graphs (PDGs) to minimally introduce parameters to the extracted method [31]. Bian et al. created extra control branches to specially handle distinct statements [9].

Meng et al. detected consistently updated clones in history and automated refactorings to remove duplicated code and to reduce repetitive coding effort [36]. Nguyen et al. leveraged historical software changes to provide API recommendations for developers [38]. Tsantalis et al. automatically detected any differences between clones, and further assessed whether it was safe to parameterize those differences without changing program behaviors [49]. Tsantalis et al. later built another tool that utilized lambda expressions to refactor clones [50].

As mentioned in prior work [36], the above-mentioned research mainly focused on the automated refactoring *feasibility* instead of *desirability*. In contrast, our research investigates the refactoring desirability. By extracting a variety of features that reflect the potential cost and benefits of refactorings to apply, we rely on machine learning to model the implicit complex interaction between features based on clones already refactored or not refactored by developers. In this way, the trained classifier can simulate the human mental process of evaluating refactoring desirability, and further suggest clones that are likely to be refactored by developers.

V. THREATS TO VALIDITY

We evaluated our approach using the clone data mined from six mature open source projects that have long version history. Our observations may not generalize well to other projects, but the trained classifier can be applied to other projects' repositories.

We did not manually inspect the non-refactored clones in the testing set, meaning that there may exist some false non-refactored clones. Please note that the false non-refactored clones in the testing set can only be caused by the true refactored clones that were missed by our automatic extraction process. However, according to existing study [51], only a

very small portion of clone groups (less than 1%) are actually refactored by developers. Therefore, the portion of false non-refactored clone groups (i.e., missing true refactored clones) in the testing set cannot exceed 1%, which would have only marginal effect on our results.

By default, CREC recommends clones for refactoring when the trained classifier predicts a refactoring likelihood greater than or equal to 50%. If a given project's version history does not have any commit that changes any clone, it is probable that none of our history-based features are helpful. Consequently, the predicted likelihood can be less than 50%, and CREC does not recommend any clone for refactoring. To at least acquire a ranked list of clones for refactoring in such scenarios, users can tune our likelihood threshold parameter to a smaller value, such as 30%.

VI. CONCLUSION

This paper presents CREC, a learning-based approach that suggests clones for *Extract Method* refactoring, no matter whether the clones are located in the same or different files.

CREC is inspired by prior work, but improves over the state-of-the-art approach by predicting refactorings based on *desirability* instead of *refactorability*. Although it is challenging to mimic developers' mental process of refactoring decisions, CREC (1) conducts static analysis to characterize the present status and past evolution of software, and (2) nicely combines static analysis with machine learning by training a classifier with the analyzed characteristics. Our evaluation evidences the effectiveness of our approach. More importantly, we observed that history-based features work more effectively than those features extracted from the current version of software. It indicates a very interesting insight: when refactoring software, developers consider more on the past evolution than the current version. By inspecting the scenarios where CREC cannot predict well, we found that the refactoring decisions seemed to be based on new software requirements or feature additions. In the future, we plan to propose approaches to better handle such scenarios. We also plan to conduct user studies with developers to obtain their opinions on the suggested clones for refactoring.

Compared with prior work, another significant improvement CREC achieves is its automatic extraction of refactored clones from version history. With our tool and data publicly available (<https://github.com/soniapku/CREC>), other researchers can access the data, reuse CREC to collect more data as needed, and further investigate refactoring recommendation and automation.

ACKNOWLEDGMENT

We thank anonymous reviewers for their thorough and valuable comments on our earlier version of the paper. This work is sponsored by the National Key Research and Development Program under Grant No. 2017YFB1001803, the National Natural Science Foundation of China under Grant No. 61672045 and 6133201, NSF Grant No. CCF-1565827, ONR Grant No. N00014-17-1-2498, NSF Grant No. CNS-1748109 and DHS Grant No. DHS-14-ST-062-001.

REFERENCES

- [1] Use JDT ASTParser to Parse Single .java files. <http://www.programcreek.com/2011/11/use-jdt-astparser-to-parse-java-file/>.
- [2] Weka 3: Data Mining Software in Java. <https://www.cs.waikato.ac.nz/ml/weka/>.
- [3] Russell, Stuart J. and Norvig, Peter. Artificial Intelligence: A Modern Approach, 2003.
- [4] L. Aversano, L. Cerulo, and M. Di Penta. How clones are maintained: An empirical study. In *Software Maintenance and Reengineering, 2007. CSMR'07. 11th European Conference on*, pages 81–90. IEEE, 2007.
- [5] M. Balazinska, E. Merlo, M. Dagenais, B. Lague, and K. Kontogiannis. Measuring clone based reengineering opportunities. In *METRICS*, page 292, 1999.
- [6] M. Balazinska, E. Merlo, M. Dagenais, B. Lague, and K. Kontogiannis. Partial redesign of java software systems based on clone analysis. In *WCRE*, page 326, 1999.
- [7] M. Balazinska, E. Merlo, M. Dagenais, B. Lague, and K. Kontogiannis. Advanced clone-analysis to support object-oriented system refactoring. In *Reverse Engineering, 2000. Proceedings. Seventh Working Conference on*, pages 98–107. IEEE, 2000.
- [8] S. Bazrafshan and R. Koschke. An empirical study of clone removals. In *IEEE International Conference on Software Maintenance*, pages 50–59, 2013.
- [9] Y. Bian, G. Koru, X. Su, and P. Ma. Spape: A semantic-preserving amorphous procedure extraction method for near-miss clones. *Journal of Systems and Software*, 86(8):2077–2093, 2013.
- [10] D. Cai and M. Kim. An empirical study of long-lived code clones. In *FASE*, volume 11, pages 432–446. Springer, 2011.
- [11] R. Caruana and A. Niculescu-Mizil. An empirical comparison of supervised learning algorithms. In *Proceedings of the 23rd international conference on Machine learning*, pages 161–168. ACM, 2006.
- [12] M. Fowler. *Refactoring: Improving the Design of Existing Code*. Addison-Wesley Professional, 2000.
- [13] Y. Freund and R. E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of computer and system sciences*, 55(1):119–139, 1997.
- [14] N. Göde. Clone removal: Fact or fiction? In *Proceedings of the 4th International Workshop on Software Clones*, pages 33–40. ACM, 2010.
- [15] N. Göde and R. Koschke. Frequency and risks of changes to clones. In *Proceedings of the 33rd International Conference on Software Engineering*, pages 311–320. ACM, 2011.
- [16] A. Goto, N. Yoshida, M. Ioka, E. Choi, and K. Inoue. How to extract differences from similar programs? a cohesion metric approach. In *7th International Workshop on Software Clones (IWSC)*, 2013.
- [17] Y. Higo, T. Kamiya, S. Kusumoto, and K. Inoue. Refactoring support based on code clone analysis. In *International Conference on Product Focused Software Process Improvement*, pages 220–233. Springer, 2004.
- [18] Y. Higo and S. Kusumoto. Identifying clone removal opportunities based on co-evolution analysis. In *Proceedings of the 2013 International Workshop on Principles of Software Evolution*, 2013.
- [19] Y. Higo, S. Kusumoto, and K. Inoue. A metric-based approach to identifying refactoring opportunities for merging code clones in a java software system. *Journal of Software: Evolution and Process*, 20(6):435–461, 2008.
- [20] T. K. Ho. Random decision forests. In *Proceedings of the Third International Conference on Document Analysis and Recognition (Volume 1) - Volume 1*, 1995.
- [21] K. Hotta, Y. Higo, and S. Kusumoto. Identifying, tailoring, and suggesting form template method refactoring opportunities with program dependence graph. *2011 15th European Conference on Software Maintenance and Reengineering*, 0:53–62, 2012.
- [22] W. Iba and P. Langley. Induction of one-level decision trees. In *Proceedings of the ninth international conference on machine learning*, pages 233–240, 1992.
- [23] L. Jiang, G. Mishnerghi, Z. Su, and S. Glondu. Deckard: Scalable and accurate tree-based detection of code clones. In *Proceedings of the 29th international conference on Software Engineering*, pages 96–105. IEEE Computer Society, 2007.
- [24] L. Jiang, Z. Su, and E. Chiu. Context-based detection of clone-related bugs. In *Proceedings of the the 6th joint meeting of the European software engineering conference and the ACM SIGSOFT symposium on The foundations of software engineering*, pages 55–64. ACM, 2007.
- [25] N. Juillerat and B. Hirsbrunner. Toward an implementation of the “form template method” refactoring. *SCAM*, 0:81–90, 2007.
- [26] T. Kamiya, S. Kusumoto, and K. Inoue. CCFinder: A multilinguistic token-based code clone detection system for large scale source code. *TSE*, pages 654–670, 2002.
- [27] M. Kim, V. Sazawal, D. Notkin, and G. Murphy. An empirical study of code clone genealogies. In *ACM SIGSOFT Software Engineering Notes*, volume 30, pages 187–196. ACM, 2005.
- [28] M. Kim, T. Zimmermann, and N. Nagappan. An empirical study of refactoring challenges and benefits at microsoft. *IEEE Transactions on Software Engineering*, 40(7):633–649, 2014.
- [29] J. Krinke. Identifying similar code with program dependence graphs. In *WCRE*, page 301, 2001.
- [30] J. Krinke. A study of consistent and inconsistent changes to code clones. In *Reverse Engineering, 2007. WCRE 2007. 14th Working Conference on*, pages 170–178. IEEE, 2007.
- [31] G. P. Krishnan and N. Tsantalis. Refactoring clones: An optimization problem. *ICSM*, 0:360–363, 2013.
- [32] V. Levenshtein. Binary Codes Capable of Correcting Deletions, Insertions and Reversals. *Soviet Physics Doklady*, 10:707, 1966.
- [33] Z. Li, S. Lu, S. Myagmar, and Y. Zhou. CP-Miner: A tool for finding copy-paste and related bugs in operating system code. pages 289–302, 2004.
- [34] Y. Lin, Z. Xing, Y. Xue, Y. Liu, X. Peng, J. Sun, and W. Zhao. Detecting differences across multiple instances of code clones. In *Proceedings of the 36th International Conference on Software Engineering*, pages 164–174. ACM, 2014.
- [35] M. Mandal, C. K. Roy, and K. A. Schneider. Automatic ranking of clones for refactoring through mining association rules. In *Software Maintenance, Reengineering and Reverse Engineering (CSMR-WCRE), 2014 Software Evolution Week-IEEE Conference on*, pages 114–123. IEEE, 2014.
- [36] N. Meng, L. Hua, M. Kim, and K. S. McKinley. Does automated refactoring obviate systematic editing? In *Software Engineering (ICSE), 2015 IEEE/ACM 37th IEEE International Conference on*, volume 1, pages 392–402. IEEE, 2015.
- [37] R. Moser, W. Pedrycz, and G. Succi. A comparative analysis of the efficiency of change metrics and static code attributes for defect prediction. In *Proceedings of the 30th International Conference on Software Engineering*, ICSE ’08, pages 181–190. New York, NY, USA, 2008. ACM.
- [38] A. T. Nguyen, M. Hilton, M. Codoban, H. A. Nguyen, L. Mast, E. Rademacher, T. N. Nguyen, and D. Dig. Api code recommendation using statistical learning from fine-grained changes. In *ACM Sigsoft International Symposium on Foundations of Software Engineering*, pages 511–522, 2016.
- [39] J. Park, M. Kim, B. Ray, and D.-H. Bae. An empirical study of supplementary bug fixes. In *IEEE Working Conference on Mining Software Repositories*, pages 40–49, 2012.
- [40] J. C. Platt. Sequential minimal optimization: A fast algorithm for training support vector machines. Technical report, ADVANCES IN KERNEL METHODS - SUPPORT VECTOR LEARNING, 1998.
- [41] J. R. Quinlan. Induction of decision trees. *Mach. Learn.*, 1986.
- [42] B. Ray, M. Kim, S. Person, and N. Rungta. Detecting and characterizing semantic inconsistencies in ported code. In *Proceedings of the 28th IEEE/ACM International Conference on Automated Software Engineering*, 2013.
- [43] C. K. Roy and J. R. Cordy. Nicad: Accurate detection of near-miss intentional clones using flexible pretty-printing and code normalization. In *Program Comprehension, 2008. ICPC 2008. The 16th IEEE International Conference on*, pages 172–181. IEEE, 2008.
- [44] H. Sajjani, V. Saini, J. Svajlenko, C. K. Roy, and C. V. Lopes. Sourcerccc: Scaling code clone detection to big-code. In *Software Engineering (ICSE), 2016 IEEE/ACM 38th International Conference on*, pages 1157–1168. IEEE, 2016.
- [45] D. Silva, N. Tsantalis, and M. T. Valente. Why we refactor? confessions of github contributors. In *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering*, pages 858–870. ACM, 2016.
- [46] R. Tairas and J. Gray. Increasing clone maintenance support by unifying clone detection and refactoring activities. *Information and Software Technology*, 54(12):1297–1307, 2012.

- [47] S. Thummalapenta, L. Cerulo, L. Aversano, and M. Di Penta. An empirical study on the maintenance of source code clones. *Empirical Software Engineering*, 15(1):1–34, 2010.
- [48] N. Tsantalis and A. Chatzigeorgiou. Ranking refactoring suggestions based on historical volatility. In *Proceedings of the 2011 15th European Conference on Software Maintenance and Reengineering*, pages 25–34, Washington, DC, USA, 2011. IEEE Computer Society.
- [49] N. Tsantalis, D. Mazinianian, and G. P. Krishnan. Assessing the refactorability of software clones. *IEEE Transactions on Software Engineering*, 41(11):1055–1090, 2015.
- [50] N. Tsantalis, D. Mazinianian, and S. Rostami. Clone refactoring with lambda expressions. In *Proceedings of the 39th International Conference on Software Engineering*, pages 60–70. IEEE Press, 2017.
- [51] W. Wang and M. W. Godfrey. Recommending clones for refactoring using design, context, and history. In *Software Maintenance and Evolution (ICSME), 2014 IEEE International Conference on*, pages 331–340. IEEE, 2014.
- [52] X. Wang, Y. Dang, L. Zhang, D. Zhang, E. Lan, and H. Mei. Predicting consistency-maintenance requirement of code clones at copy-and-paste time. *IEEE Transactions on Software Engineering*, 40(8):773–794, Aug 2014.